

zetesis-puremath Blueprint

A Lean 4 pure-mathematics reservoir:
analytic measurability, information theory, concentration, minimax,
kernel embeddings, and dual VC.

Dhruv Gupta

April 2026

Contents

1	Introduction	1
1.1	What is formalized here	1
1.2	Verification tiers	1
2	Measure theory	3
2.1	Analytic measurability	3
2.2	Choquet capacity	3
2.3	Baire-space cylinder infrastructure	3
2.4	Total variation distance	4
3	Information theory	5
3.1	Real-valued KL bridge	5
3.2	Binary KL and Pinsker's inequality	5
3.3	Mutual information	6
4	Probability	7
4.1	Bounded random variables + concentration	7
4.2	Exchangeability and double-sample	7
4.3	Finite-type PMF	7
5	Decision theory: minimax and MWU	9
5.1	Boolean-matrix zero-sum games	9
5.2	Multiplicative weights update	9
5.3	Approximate minimax	10
6	Reproducing-kernel Hilbert spaces	11
6.1	Bounded kernels	11
6.2	Kernel mean embedding	11
6.3	Maximum Mean Discrepancy	11
6.4	Hilbert-Schmidt Independence Criterion	12
7	Combinatorics: dual VC	13
7.1	Binary matrix setup	13
7.2	Sauer-Shelah corollary	13
7.3	Assouad's dual VC bound	13
	References	15

Chapter 1

Introduction

This blueprint pins every mathematical definition and headline theorem in `zetesis-puremath`¹ to its Lean 4 declaration. The companion source is a personal pure-math reservoir: content needed by the author’s formalization projects (`formal-learning-theory-kernel`, SMFE, and others) that is either absent from Mathlib or specialized in a way Mathlib’s current formulation does not support directly.

1.1 What is formalized here

Six substrates, each with its own chapter below:

- **Measure theory on Polish spaces.** Analytic sets are null-measurable under finite Borel measures; Kechris 30.13 capacitability theorem for analytic sets; real-valued metric total variation distance.
- **Information theory.** Real-valued Kullback-Leibler bridge to Mathlib’s `ENNReal`-valued `klDiv`; binary KL with a full derivative / monotonicity chain; Pinsker’s inequality with the sharp constant; mutual information and its zero-iff-product characterization.
- **Probability.** `BoundedRandomVariable` typeclass with a Chebyshev-via-majority bound; double-sample / `ValidSplit` / `SplitMeasure` exchangeability infrastructure; `FintypePMF` (an $\mathbb{R}$ -valued, `linarith`-compatible specialization of Mathlib’s `PMF`).
- **Decision theory.** Covering-based minimax; multiplicative-weights-update (MWU) in the Freund-Schapire $\beta = (1 - \eta)$ convention; approximate-minimax bound $T = \mathcal{O}(\log N / \varepsilon^2)$.
- **Reproducing-kernel Hilbert spaces.** Kernel mean embedding with the reproducing property; Maximum Mean Discrepancy (MMD) and characteristic kernels; Hilbert-Schmidt Independence Criterion (HSIC).
- **Combinatorics.** Assouad’s 1983 dual VC bound $\text{VCdim}(M^\top) \leq 2^{d+1} - 1$.

1.2 Verification tiers

All 114 public theorems discharge under the following tiers, all green at commit 471e5d5 (see [test/verification/CI/](#)):

- **Tier 0:** lake build — 0 errors, 0 warnings, 0 sorry in ZPM/.

¹<https://github.com/Zetetic-Dhruv/zetesis-puremath>

- **Tier 1:** `#print axioms` — 113 theorems use only `propext`, `Classical.choice`, `Quot.sound`; the 114th (`challenge_bnd_subset_cyl`) uses none.
- **Tier 2:** `leanchecker -fresh` per module (GitHub Actions `workflow_dispatch`).
- **Tier 3:** `comparator` + Landlock sandbox on GitLab CI; “*Lean default kernel accepts the solution. Your solution is okay!*” (pipeline #2463996501, duration 36 min).

Every statement in this blueprint carries a `\lean{...}` annotation linking to the Lean declaration; where the full proof is in the source, `\leanok` marks the entry as formally verified.

Chapter 2

Measure theory

2.1 Analytic measurability

Theorem 2.1 (Analytic null-measurability). *Lean: MeasureTheory.AnalyticSet.nullMeasurableSet*
Let α be a Polish space, μ a finite Borel measure on α , and $S \subseteq \alpha$ analytic. Then S is μ -null-measurable.

Theorem 2.2 (Inner approximation by compacts). *Lean: MeasureTheory.AnalyticSet.exists_isCompact_measureReal*
Under the hypotheses of Theorem 2.1, for every $\delta > 0$ there exists a compact $K \subseteq S$ with $\mu(S) - \mu(K) < \delta$ (in $\mathbb{R}$ -valued form).

2.2 Choquet capacity

Definition 2.3 (Choquet capacity). *Lean: MeasureTheory.IsChoquetCapacity*

A map $I : \text{Set}(\alpha) \rightarrow \mathbb{R}_{\geq 0}^{\infty}$ is a *Choquet capacity* iff (i) it is monotone, (ii) continuous from below along increasing sequences, and (iii) continuous from above along antitone sequences of closed sets. Every finite Borel measure on a Polish space is a Choquet capacity [8].

Definition 2.4 (Compact capacity hull). *Lean: MeasureTheory.compactCap*

For μ and S as above,

$$\text{compactCap } \mu S := \sup\{\mu(K) : K \subseteq S, K \text{ compact}\}.$$

Theorem 2.5 (Kechris 30.13: analytic capacitability). *Lean: MeasureTheory.AnalyticSet.cap_eq_iSup_isCompact*
Let I be a Choquet capacity on a Polish space α , and S an analytic subset of α . Then

$$I(S) = \sup\{I(K) : K \subseteq S, K \text{ compact}\}.$$

The proof is the Kechris 30.13 diagonal construction on cylinders in Baire space $\mathbb{N}^{\mathbb{N}}$ [14]. Together with Definition 2.4 and the fact that every finite Borel measure is a Choquet capacity, this specializes to $\mu(S) = \text{compactCap } \mu S$ (Theorem 2.1 is an immediate corollary).

2.3 Baire-space cylinder infrastructure

The proof of Theorem 2.5 requires a small theory of cylinders in $\mathbb{N}^{\mathbb{N}}$.

Definition 2.6 (Cylinder). *Lean: Cyl*

For $N : \mathbb{N} \rightarrow \mathbb{N}$ and $n \in \mathbb{N}$,

$$\text{Cyl } N n := \{g : \mathbb{N} \rightarrow \mathbb{N} \mid \forall i \leq n, g i \leq N i\}.$$

Definition 2.7 (Bounded box). *Lean: Bnd*

$$\text{Bnd } N := \bigcap_n \text{Cyl } N \, n = \prod_i [0, N_i].$$

Definition 2.8 (Truncation map). *Lean: truncate*

$\text{truncate } N \, g \, i := \min(g \, i, N \, i)$. Sends any sequence into $\text{Bnd } N$ and agrees with g on cylinders.

Lemma 2.9 ($\text{Bnd } N$ is compact). *Lean: isCompact_bnd*

$\text{Bnd } N$ is compact (Tychonoff applied to the finite intervals $[0, N_i]$).

Lemma 2.10 ($\text{Bnd} \subseteq \text{Cyl}$). *Lean: bnd_subset_cyl*

$\text{Bnd } N \subseteq \text{Cyl } N \, n$ for every n .

Lemma 2.11 (Cylinder successor splits along the next coordinate). *Lean: cyl_succ_eq*

$$\text{Cyl } N \, (n + 1) = \text{Cyl } N \, n \cap \{g : g \, (n + 1) \leq N \, (n + 1)\}.$$

Lemma 2.12 (Intersection with a coordinate update). *Lean: cyl_inter_eq_cyl_update*

Intersecting a cylinder with a coordinate-level constraint rewrites as a cylinder with an updated bound at that coordinate.

Lemma 2.13 (Monotone split). *Lean: monotone_cyl_split*

The map $k \mapsto \text{Cyl } N \, n \cap \{g : g \, (n + 1) \leq k\}$ is monotone in k .

Lemma 2.14 (Truncate lands in Bnd). *Lean: truncate_mem_bnd*

$\text{truncate } N \, g \in \text{Bnd } N$ for every g .

Lemma 2.15 (Truncate agrees on cylinders). *Lean: truncate_agree_on_cyl*

If $g \in \text{Cyl } N \, n$ then $\text{truncate } N \, g \, i = g \, i$ for all $i \leq n$.

Lemma 2.16 (Cylinder extensionality). *Lean: cyl_ext*

Two elements of $\text{Cyl } N \, n$ agreeing on the first $n + 1$ coordinates are equal on Cyl .

Lemma 2.17 (Intersection of cylinder-closure images equals compact image). *Lean: iInter_closure_image_cyl_eq*

Under a continuous parametrization $f : \mathbb{N}^\mathbb{N} \rightarrow \alpha$ into a Polish space, $\bigcap_n \overline{f(\text{Cyl } N \, n)} = f(\text{Bnd } N)$.

2.4 Total variation distance

Definition 2.18 (Real-valued total variation). *Lean: MeasureTheory.tvDistReal*

For finite Borel measures P, Q on α ,

$$\text{TV}(P, Q) := \sup_A |P(A).toReal - Q(A).toReal|$$

where the supremum is over measurable $A \subseteq \alpha$. This is the *half-sup* / metric convention (Billingsley, Villani), chosen so that Pinsker reads as $\text{TV}(P, Q) \leq \sqrt{\text{KL}(P \| Q)}/2$. The lemmas `tvDistReal_set_nonempty` and `tvDistReal_set_bddAbove` discharge the nonemptiness and boundedness conditions that $\mathbb{R}$ -valued `sSup` requires.

Remark 2.19. The repository also contains a second total-variation object, `FintypePMF.tvDistance` $p \, q = \sum_a |p \, a - q \, a|$ (the L^1 convention, see Chapter 4); the two conventions differ by a factor of 2 on discrete supports. Downstream consumers should not compose them without inserting the factor-of- $\frac{1}{2}$ correction.

Chapter 3

Information theory

3.1 Real-valued KL bridge

Definition 3.1 (Real-valued KL). *Lean: `InformationTheory.klDivReal`*

For measures P, Q on a measurable space, set

$$\text{KL}_{\mathbb{R}}(P, Q) := \begin{cases} \int \log \frac{dP}{dQ} dP & \text{if } P \ll Q, \\ 0 & \text{otherwise.} \end{cases}$$

This bridges Mathlib’s `ENNReal`-valued `klDiv` to $\mathbb{R}$ for `linarith`-compatible arithmetic; the identity $\text{KL}_{\mathbb{R}}(P, Q) = (\text{KL}(P, Q)).\text{toReal}$ holds unconditionally (`klDivReal_eq_toReal_klDiv`).

Theorem 3.2 (Non-negativity). *Lean: `InformationTheory.klDivReal_nonneg`*

$\text{KL}_{\mathbb{R}}(P, Q) \geq 0$ for all measures (Gibbs’ inequality).

Theorem 3.3 (Zero iff equal, under finiteness). *Lean: `InformationTheory.klDivReal_eq_zero_iff`*

If $P \ll Q$, $\text{KL}(P, Q) \neq \infty$, and both are finite measures, then $\text{KL}_{\mathbb{R}}(P, Q) = 0 \iff P = Q$ [7, Thm 2.6.3].

3.2 Binary KL and Pinsker’s inequality

Definition 3.4 (Binary KL). *Lean: `InformationTheory.klBin`*

For $p, q \in [0, 1]$,

$$\text{KL}_{\text{bin}}(p, q) := p \log \frac{p}{q} + (1 - p) \log \frac{1-p}{1-q}$$

with the convention $0 \cdot \log(0/\cdot) = 0$.

Lemma 3.5 (Derivative in q). *Lean: `InformationTheory.hasDerivAt_klBin_q`*

For $q \in (0, 1)$, $\partial_q \text{KL}_{\text{bin}}(p, q) = (q - p)/(q(1 - q))$.

Theorem 3.6 (Binary Pinsker, sharp constant 2). *Lean: `InformationTheory.binary_pinsker`*

For $p \in [0, 1]$ and $q \in (0, 1)$, $2(p - q)^2 \leq \text{KL}_{\text{bin}}(p, q)$.

Theorem 3.7 (Pinsker’s inequality). *Lean: `InformationTheory.pinsker_proof`*

For probability measures $P \ll Q$ with $\text{KL}(P, Q) < \infty$,

$$\text{TV}(P, Q) \leq \sqrt{\text{KL}_{\mathbb{R}}(P, Q)/2}$$

[7, Thm 11.6.1]. The convention matches Definition 2.18 (half-sup TV); with the L^1 convention the bound reads $\|P - Q\|_1 \leq \sqrt{2 \text{KL}}$.

Theorem 3.8 (Binary-to-general coarsening). *Lean: `InformationTheory.klBin_le_klDivReal`*

For every measurable A with $P(A) = p$, $Q(A) = q$, $\text{KL}_{\text{bin}}(p, q) \leq \text{KL}_{\mathbb{R}}(P, Q)$. This is the two-point data-processing inequality and is the step linking the binary Pinsker to the general measure-theoretic Pinsker.

3.3 Mutual information

Definition 3.9 (Real-valued mutual information). *Lean: `InformationTheory.mutualInformationReal`*

For a probability measure P on $X \times Y$ with marginals P_X, P_Y , $\text{MI}_{\mathbb{R}}(P) := \text{KL}_{\mathbb{R}}(P, P_X \otimes P_Y)$.

Theorem 3.10 (Non-negativity of mutual information). *Lean: `InformationTheory.mutualInformationReal_nonneg`*

$\text{MI}_{\mathbb{R}}(P) \geq 0$ (Gibbs applied to the joint vs. product-of-marginals).

Theorem 3.11 (Mutual information vanishes iff independent). *Lean: `InformationTheory.mutualInformationReal_eq`*

Under $P \ll P_X \otimes P_Y$ and $\text{KL}(P, P_X \otimes P_Y) \neq \infty$, $\text{MI}_{\mathbb{R}}(P) = 0 \iff P = P_X \otimes P_Y$.

Chapter 4

Probability

4.1 Bounded random variables + concentration

Definition 4.1 (Bounded random variable). `Lean: ProbabilityTheory.BoundedRandomVariable`

$f : \Omega \rightarrow \mathbb{R}$ is *bounded* in $[a, b]$ under μ iff f is measurable and $f(\omega) \in [a, b]$ for μ -almost every ω .

Theorem 4.2 (Chebyshev majority bound). `Lean: ProbabilityTheory.chebyshev_majority_bound`

Let $X_1, \dots, X_k$ be independent random variables with each $\mu(A_j) \geq 2/3$ where $A_j = \{X_j = \text{correct}\}$. Then for $k \geq 9/\delta$, the probability that a majority of the X_j is correct is at least $1 - \delta$.
Proof via Popoviciu's variance bound + Chebyshev's inequality [5, §2].

4.2 Exchangeability and double-sample

Definition 4.3 (Sample bundle). `Lean: ProbabilityTheory.ExchangeableSample`

A bundle $\langle m, \mu, \text{IsProbabilityMeasure } \mu \rangle$ where $m \in \mathbb{N}$ is the sample size and μ is a probability measure on the sample space. Used with the product measure $\mu^{\otimes 2m}$ as the natural exchangeable double-sample domain.

Definition 4.4 (Valid split). `Lean: ProbabilityTheory.ValidSplit`

$s : \text{Fin}(2m) \rightarrow \text{Bool}$ is a *valid split* iff exactly m of its values are `true`. There are $\binom{2m}{m}$ valid splits (Vapnik-Chervonenkis double-sample [23]).

Definition 4.5 (Uniform split measure). `Lean: ProbabilityTheory.SplitMeasure`

The uniform probability measure on the set of valid splits. This is the symmetrization measure for the double-sample trick.

4.3 Finite-type PMF

Definition 4.6 (Real-valued fintype PMF). `Lean: ProbabilityTheory.FintypePMF`

For a fintype α , a `FintypePMF` α is a pair $\langle p, h_{\geq 0}, h_{\sum=1} \rangle$ where $p : \alpha \rightarrow \mathbb{R}$, $p \geq 0$ pointwise, and $\sum_{a:\alpha} p(a) = 1$. This is an $\mathbb{R}$ -valued, `linarith`-compatible specialization of Mathlib's $\mathbb{R}_{\geq 0}^\infty$ -valued PMF [5, §4].

Theorem 4.7 (Bridge to Mathlib PMF). `Lean: ProbabilityTheory.FintypePMF.toPMF_apply`

The map `toPMF` : `FintypePMF` $\alpha \rightarrow \text{PMF } \alpha$ satisfies $(\text{toPMF } p)(a) = \text{ENNReal.ofReal}(p.\text{prob } a)$, which is lossless on valid `FintypePMF` inputs.

Definition 4.8 (L^1 TV distance). *Lean: ProbabilityTheory.FintypePMF.tvDistance*

$\text{tvDistance } p \, q := \sum_{a:\alpha} |p(a) - q(a)|$. This is the *full* L^1 convention, not the half-sup TV of Definition 2.18; on discrete supports they differ by a factor of 2.

Theorem 4.9 (Non-negativity of tvDistance). *Lean: ProbabilityTheory.FintypePMF.tvDistance_nonneg*
 $\text{tvDistance}(p, q) \geq 0$.

Theorem 4.10 (Symmetry of tvDistance). *Lean: ProbabilityTheory.FintypePMF.tvDistance_comm*
 $\text{tvDistance}(p, q) = \text{tvDistance}(q, p)$.

Theorem 4.11 (Self-distance is zero). *Lean: ProbabilityTheory.FintypePMF.tvDistance_self*
 $\text{tvDistance}(p, p) = 0$.

Theorem 4.12 (Transfer principle). *Lean: ProbabilityTheory.FintypePMF.expectation_approx_of_tv*
 For Boolean-valued f , $|\mathbb{E}_p[f] - \mathbb{E}_q[f]| \leq \text{tvDistance}(p, q)$. In the half-sup TV convention this would read $\leq 2 \cdot \text{TV}(p, q)$.

Chapter 5

Decision theory: minimax and MWU

5.1 Boolean-matrix zero-sum games

Definition 5.1 (Boolean-matrix payoff). *Lean: ProbabilityTheory.boolGamePayoff*

For a Boolean matrix $M : R \rightarrow C \rightarrow \text{Bool}$, a row distribution $p : \text{FintypePMF } R$, and a pure column $c : C$,

$$\text{boolGamePayoff } M \, p \, c := \sum_{r:R} p(r) \cdot [M \, r \, c].$$

This is the expected $\{0,1\}$ -gain of the row player, against a pure column play, in the matrix game M [10].

Theorem 5.2 (One-sided minimax sanity). *Lean: ProbabilityTheory.minimax_value_le_one*

If the row player can guarantee value v against every mixed column (the *hrow hypothesis*), then $v \leq 1$.

5.2 Multiplicative weights update

Definition 5.3 (MWU config). *Lean: ProbabilityTheory.MWUConfig*

An $\text{MWUConfig } C$ is a pair $\langle w, h_{>0} \rangle$ where $w : C \rightarrow \mathbb{R}$ is the strictly positive weight vector. The *potential* is $\Phi(w) := \sum_c w(c)$, the *induced PMF* is $\text{toPMF}(w)(c) := w(c)/\Phi(w)$ [3, Thm 2.1].

Definition 5.4 (Update and closed-form). *Lean: ProbabilityTheory.mwuUpdateWeights*

Given a row $r : R$ and step size $\eta < 1$, the update is $w'(c) := w(c) \cdot (1 - \eta)$ if $M \, r \, c = \text{true}$, else $w'(c) := w(c)$. This is the Freund-Schapire $\beta = (1 - \eta)$ Hedge convention [9]. The closed form

$$w_T(c) = w_0(c) \cdot (1 - \eta)^{\text{hitCount}_T(c)}$$

is Theorem `mwu_weight_eq_pow_hitCount`.

Theorem 5.5 (One-step potential bound). *Lean: ProbabilityTheory.potential_one_step_bound*

If the row player can guarantee value v against every mixed column and $0 \leq \eta < 1$, then after one MWU step $\Phi_{t+1} \leq \Phi_t \cdot (1 - \eta v)$.

Theorem 5.6 (T -step potential bound). *Lean: ProbabilityTheory.mwu_potential_T_bound*

Induction gives $\Phi_T \leq |C| \cdot (1 - \eta v)^T$.

5.3 Approximate minimax

Theorem 5.7 (Approximate minimax via MWU). *Lean: `ProbabilityTheory.mwu_approx_minimax`*

Fix $\varepsilon > 0$, set $\eta = \varepsilon/4$ and $T = \max\{1, \lceil 16 \log N / \varepsilon^2 \rceil\}$ where $N = |C|$. If the row player can guarantee value v against every mixed column, then the empirical row mixture produced by T rounds of MWU achieves

$$v - \varepsilon \leq \text{boolGamePayoff } M \text{ } p \text{ } c \quad \text{for every pure } c \in C.$$

This is the Freund-Schapire constructive minimax [10, Thm 3]; [3, §3.1] unifies this with general online learning.

Theorem 5.8 (Covering-based minimax existence). *Lean: `ProbabilityTheory.covering_minimax`*

Given the `hrow` hypothesis, there exists a mixed-row PMF whose payoff against every pure column is at least $1/|C|$. Proof by pigeonhole-averaging over per-column pure-row witnesses extracted via `ProbabilityTheory.exists_covering_row`.

Theorem 5.9 (ε -approximate finite minimax). *Lean: `ProbabilityTheory.finite_approx_minimax`*

Under a feasibility side-condition $v - \varepsilon \leq 1/|C|$, the covering-based mixed-row strategy from Theorem 5.8 achieves payoff $\geq v - \varepsilon$ against every pure column. This is the logical dual of the quantitative MWU bound in Theorem 5.7.

Chapter 6

Reproducing-kernel Hilbert spaces

Background conventions: the Mathlib class `RKHS` `H X` is the scalar-valued reproducing-kernel Hilbert space on a domain X , with Hilbert space $\mathcal{H} = H$ and feature map $k_x := \text{kerFun } H \ x \ 1 \in H$ satisfying $f(x) = \langle k_x, f \rangle$ [1]. Throughout, $\mathcal{H}$ denotes a complete, separable, real inner product space.

6.1 Bounded kernels

Definition 6.1 (Bounded kernel). *Lean: `BoundedKernel`*

A `BoundedKernel` $\mathcal{H}$ is a witness $\langle C, h_{\geq 0}, h_{\leq} \rangle$ with $C \geq 0$ and $\|k_x\| \leq C$ for all $x : X$. Bounded-feature kernels are the canonical setting for kernel mean embeddings [21, §4.2].

Theorem 6.2 (Evaluation bound). *Lean: `norm_apply_le_bound_mul_norm`*

$|f(x)| \leq C \cdot \|f\|$ for every $f \in \mathcal{H}$ and $x \in X$.

6.2 Kernel mean embedding

Definition 6.3 (Kernel mean embedding). *Lean: `kernelMeanEmbedding`*

For a measure P on X , set

$$\mu_P := \int_X k_x \, dP(x) \in \mathcal{H}$$

(Bochner integral) [19].

Theorem 6.4 (Reproducing property for the embedding). *Lean: `kernelMeanEmbedding_reproducing`*

For integrable f , $\langle \mu_P, f \rangle = \int_X f(x) \, dP(x)$ [11, §2, eq. (2)].

Theorem 6.5 (Norm bound). *Lean: `kernelMeanEmbedding_norm_le`*

Under *IsProbabilityMeasure* P , $\|\mu_P\| \leq C$.

6.3 Maximum Mean Discrepancy

Definition 6.6 (Squared MMD). *Lean: `mmdSq`*

$\text{MMD}^2(P, Q) := \|\mu_P - \mu_Q\|^2$ [11, Defn. 2].

Theorem 6.7 (Non-negativity of MMD^2). *Lean: `mmdSq_nonneg`*

$\text{MMD}^2(P, Q) \geq 0$.

Theorem 6.8 (MMD² on identical measures). *Lean: mmdSq_self*

$$\text{MMD}^2(P, P) = 0.$$

Theorem 6.9 (Symmetry of MMD²). *Lean: mmdSq_comm*

$$\text{MMD}^2(P, Q) = \text{MMD}^2(Q, P).$$

Definition 6.10 (Characteristic RKHS). *Lean: IsCharacteristic*

$\mathcal{H}$ is *characteristic on X* iff the kernel mean embedding $P \mapsto \mu_P$ is injective. This is the modern definition [20, Defn. 6].

Theorem 6.11 (Zero-iff-equal under characteristic kernel). *Lean: mmdSq_zero_iff*

$$\text{Under } \text{IsCharacteristic } \mathcal{H}, \text{MMD}^2(P, Q) = 0 \iff P = Q.$$

6.4 Hilbert-Schmidt Independence Criterion

Definition 6.12 (HSIC via product-space MMD). *Lean: hsicDef*

For a joint measure P on $X \times Y$ with marginals P_X, P_Y ,

$$\text{HSIC}(P) := \text{MMD}^2(P, P_X \otimes P_Y).$$

This is the modern unified formulation [20, Thm 24]; originally introduced as a cross-covariance-operator Hilbert-Schmidt norm [12].

Theorem 6.13 (Non-negativity of HSIC). *Lean: hsicDef_nonneg*

$$\text{HSIC}(P) \geq 0.$$

Theorem 6.14 (HSIC vanishes on independent measures). *Lean: hsicDef_eq_zero_of_independent*

If $P = P_X \otimes P_Y$ then $\text{HSIC}(P) = 0$ (direct consequence of $\text{MMD}^2(P, P) = 0$).

Theorem 6.15 (HSIC zero iff independent). *Lean: hsicDef_zero_iff_independent*

Under $\text{IsCharacteristic } \mathcal{H}$ on $X \times Y$, $\text{HSIC}(P) = 0 \iff P = P_X \otimes P_Y$.

Chapter 7

Combinatorics: dual VC

7.1 Binary matrix setup

Definition 7.1 (Binary matrix). *Lean*: `BinaryMatrix`

`BinaryMatrix m n` := `Fin m → Fin n → Bool`. Type-definitionally equal to `Matrix (Fin m) (Fin n) Bool`.

Definition 7.2 (Induced set family). *Lean*: `BinaryMatrix.toFinsetFamily`

The *induced set family* of a binary matrix $M : \text{Fin } m \rightarrow \text{Fin } n \rightarrow \text{Bool}$ consists of the true-supports of its rows, $\{ \{j : M \ i \ j = \text{true}\} : i \in \text{Fin } m \}$.

Definition 7.3 (Matrix shattering). *Lean*: `BinaryMatrix.shatters`

For a binary matrix M and $S \subseteq \text{Fin } n$, M *shatters* S iff every pattern $t \subseteq S$ is realized by some row of M (Vapnik–Chervonenkis 1971; Sauer 1972).

Theorem 7.4 (Consistency with `Finset.Shatters`). *Lean*: `BinaryMatrix.shatters_iff`

M *shatters* $S \iff \text{toFinsetFamily } M \text{ shatters } S \text{ in Mathlib's sense.}$

7.2 Sauer-Shelah corollary

Theorem 7.5 (Cardinality of the induced family). *Lean*: `BinaryMatrix.card_toFinsetFamily_le`

If $\text{VCdim}(\text{toFinsetFamily } M) \leq d$ then $|\text{toFinsetFamily } M| \leq \sum_{k=0}^d \binom{n}{k}$. This is the Sauer-Shelah bound [17, 23] routed through *Mathlib's* `Finset.card_shatterer_le_sum_vcDim` (Pajor + Sauer-Shelah) into the matrix setting.

7.3 Assouad's dual VC bound

Lemma 7.6 (Dual-to-primal shattering coding). *Lean*: `BinaryMatrix.transpose_shatters_imp_shatters`

If $M^\top$ *shatters* $S \subseteq \text{Fin } m$ with $|S| \geq 2^{d+1}$, then M *shatters* some $T \subseteq \text{Fin } n$ with $|T| = d + 1$. *Proof by a bitstring-coding argument: over $(d+1)$ -bit patterns on S , each pattern must be realized by a distinct column, giving $d + 1$ distinguishing columns that M shatters [2, Thm 2.13].*

Theorem 7.7 (Assouad's bound: $\text{VCdim}(M^\top) \leq 2^{d+1} - 1$). *Lean*: `BinaryMatrix.assouad_transpose_vcDim`

If $\text{VCdim}(M) \leq d$ then $\text{VCdim}(M^\top) \leq 2^{d+1} - 1$ [2, Thm 2.13]. The bound is tight at small d ; no strengthening is known for general d without additional structure.

References

Bibliography

- [1] N. Aronszajn. *Theory of reproducing kernels*. Trans. Amer. Math. Soc., 68:337–404, 1950.
- [2] P. Assouad. *Densité et dimension*. Ann. Inst. Fourier, 33(3):233–282, 1983.
- [3] S. Arora, E. Hazan, and S. Kale. *The Multiplicative Weights Update Method: A Meta-Algorithm and Applications*. Theory of Computing, 8:121–164, 2012.
- [4] V. I. Bogachev. *Measure Theory*, Vol. II. Springer, 2007.
- [5] S. Boucheron, G. Lugosi, and P. Massart. *Concentration Inequalities: A Nonasymptotic Theory of Independence*. Oxford University Press, 2013.
- [6] G. Choquet. *Theory of capacities*. Ann. Inst. Fourier, 5:131–295, 1955.
- [7] T. M. Cover and J. A. Thomas. *Elements of Information Theory*, 2nd ed. Wiley-Interscience, 2006.
- [8] C. Dellacherie and P.-A. Meyer. *Probabilities and Potential A*. North-Holland, 1978.
- [9] Y. Freund and R. E. Schapire. *A Decision-Theoretic Generalization of On-Line Learning and an Application to Boosting*. J. Comput. System Sci., 55(1):119–139, 1997.
- [10] Y. Freund and R. E. Schapire. *Adaptive Game Playing Using Multiplicative Weights*. Games Econ. Behav., 29:79–103, 1999.
- [11] A. Gretton, K. M. Borgwardt, M. J. Rasch, B. Schölkopf, and A. Smola. *A Kernel Two-Sample Test*. J. Mach. Learn. Res., 13:723–773, 2012.
- [12] A. Gretton, O. Bousquet, A. Smola, and B. Schölkopf. *Measuring Statistical Dependence with Hilbert–Schmidt Norms*. In ALT, 2005.
- [13] O. Kallenberg. *Foundations of Modern Probability*, 3rd ed. Springer, 2021.
- [14] A. S. Kechris. *Classical Descriptive Set Theory*. Graduate Texts in Mathematics, Vol. 156. Springer, 1995.
- [15] S. Kullback. *A Lower Bound for Discrimination Information in Terms of Variation*. IEEE Trans. Inform. Theory, 13:126–127, 1967.
- [16] Y. Polyanskiy and Y. Wu. *Information Theory: From Coding to Learning*. Cambridge University Press, 2025.
- [17] N. Sauer. *On the density of families of sets*. J. Combin. Theory Ser. A, 13:145–147, 1972.
- [18] S. Shelah. *A combinatorial problem; stability and order for models and theories in infinitary languages*. Pacific J. Math., 41:247–261, 1972.

- [19] A. Smola, A. Gretton, L. Song, and B. Schölkopf. *A Hilbert Space Embedding for Distributions*. In ALT, 2007.
- [20] B. K. Sriperumbudur, A. Gretton, K. Fukumizu, B. Schölkopf, and G. R. G. Lanckriet. *Hilbert Space Embeddings and Metrics on Probability Measures*. J. Mach. Learn. Res., 11:1517–1561, 2010.
- [21] I. Steinwart and A. Christmann. *Support Vector Machines*. Springer, 2008.
- [22] A. B. Tsybakov. *Introduction to Nonparametric Estimation*. Springer, 2009.
- [23] V. N. Vapnik and A. Ya. Chervonenkis. *On the Uniform Convergence of Relative Frequencies of Events to Their Probabilities*. Theory Probab. Appl., 16:264–280, 1971.
- [24] C. Villani. *Optimal Transport: Old and New*. Springer, 2009.